



Putting Queens in Carry Chains – No. 27 –

Dr.-Ing. Thomas B. Preußner
thomas.preusser@tu-dresden.de






Itinerary

- Problem and Complexity Overview
- Solution Approach
- Hardware Mapping and Optimization
- Ongoing Computation for $N = 27$

The N -Queens Puzzle

- Place N *non-attacking* Queens on a $N \times N$ chessboard:



- Generic solution templates?
→ How many (fundamental) solutions?

Motivation

The exploration of an N -Queens Puzzle is:

- an embarrassingly parallel,
- easily scalable,
- computation-bounded

workload.

It serves as:

- a training object for working with cramped designs:
 - efficient coding and resource utilization,
 - tooling and parameter exploration.
- a tool and device benchmark.

And yes: *We just can!*

Known Solution Counts

N	Solutions	N	Solutions
1	1	14	365596
2	0	15	2279184
3	0	16	14772512
4	2	17	95815104
5	10	18	666090624
6	4	19	4968057848
7	40	20	39029188884
8	92	21	314666222712
9	352	22	2691008701644
10	724	23	24233937684440
11	2680	24	227514171973736
12	14200	25	2207893435808352
13	73712	26	22317699616364044

Exhaustive backtracking solution exploration requires factorial time $O(N!)$.

Very hard beyond $N = 20$.

$N = 25$:

- Java grid computation by INRIA, France.
- Runtime:
 - Real >6 Months
 - Sequential >53 Years

Known Solution Counts

<u>N</u>	<u>Solutions</u>	<u>N</u>	<u>Solutions</u>
1	1	14	365596
2	0	15	2279184
3	0	16	14772512
4	2	17	95815104
5	10	18	666090624
6	4	19	4968057848
7	40	20	39029188884
8	92	21	314666222712
9	352	22	2691008701644
10	724	23	24233937684440
11	2680	24	227514171973736
12	14200	25	2207893435808352
13	73712	26	22317699616364044

Exhaustive backtracking solution exploration requires factorial time $O(N!)$.

Very hard beyond $N = 20$.

$N = 26$:


- 9-month computation on FPGAs completing July 11, 2009.
- Result confirmed by Russian MC# super computing project on August 30, 2009.

Tackling $N = 26$

- Embarrassingly parallel workload:
 1. Preplace $L \ll N$ columns.
 2. Explore subboards **independently**.
 3. Collect and add up subtotals.
- Ideally suited for distributed computation:
 - Internet (BOINC) → NQueens@Home
 - FPGA! → Queens@TUD
 - Challenge the power of a world-wide distributed computation effort by an intelligent FPGA implementation.
 - Identified and reported an overflow bug on November 7, 2008.
 - Thereby, resolved an open dispute on the solution for $N = 24$ *without* any own computation.

Algorithmic Overview

Exhaustive backtracking solution exploration.



valid placement yet to explore?

Yes


1. Mark *explored*.
2. Update *blocking vectors*.
3. Advance to next column / Count solution.

No

1. Clear markings.
2. Retreat to previous column / Done.

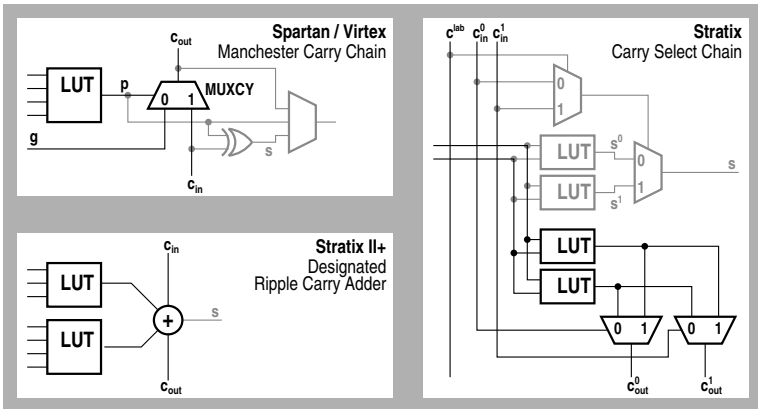
Computing *Blocking Vectors* avoids frequent constraint validation.

FPGA Mapping



Using carry chains to process one column in one *fast* clock cycle.

Carry Chain Structures



Carry chains are implemented to speed up binary word *addition*.

Generic Carry-Chain Mapping through Addition

1. Derive Carry / Token Propagation

Case	c_{i+1}	Description
k_i : Kill	0	never a carry: holding no queen and not blocked
p_i : Propagate	c_i	pass a carry: holding no queen but blocked
g_i : Generate	1	always a carry: holding current queen placement

2. Determine Addends

$$\begin{aligned} a_i &= g_i + p_i \\ b_i &= g_i \end{aligned}$$

3. Infer Token from Sum $s \leq a + b$

In equations dependent on the incoming carry/token use:

$$c_i = s_i \oplus p_i$$

Shown mapping to Xilinx devices uses optimized implementation.

Pushing Performance

Optimization for a small size and a high clock frequency:


- Maintain a single active column.
- Keep placed columns within a plain array of shifted registers.
- Use global blocking vectors for all rows and diagonals setting and unsetting placements and retreats, respectively.
Note that this is quite expensive in software!

Designing $N = 27$


Column-based pre-placement may exploit line symmetry to cut search space in half.

Designing $N = 27$

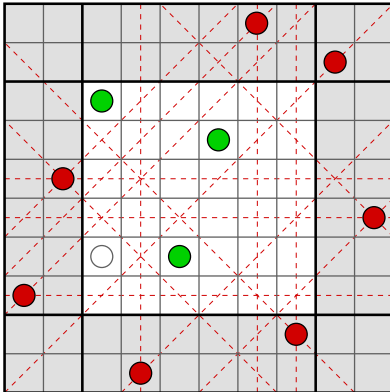
Column-based pre-placement may exploit line symmetry to cut search space in half.
There is more:




Coronal Pre-Placement



Coronal Pre-Placement



Coronal Pre-Placement



Advantage:

- Search space reduced to an eighth.









Challenges (solved):

- Define *canonical* representative.
- Count solutions of self-symmetric pre-placements correctly.

2.024.110.796 coronal pre-placements for $N = 27$.

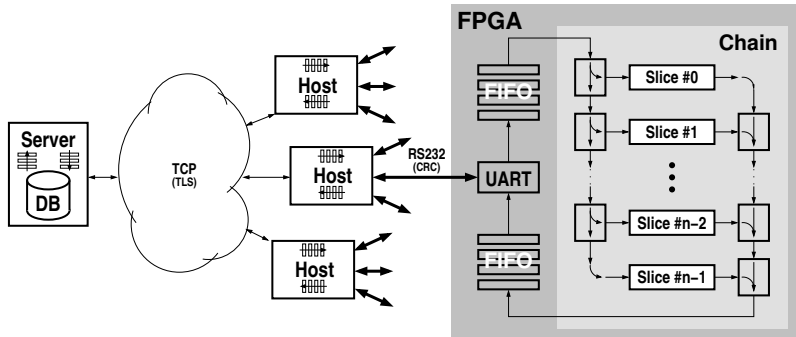
Solve one per second: 64 years of sequential computation time.

Pre-Placement: Canonical Representative



Minimum as determined by lexicographic order of *traits*.

Project Infrastructure



Scalability


Inherently computation bounded: subproblem solution is encoded in 21 bytes only.
Current peaks at 25 solutions per second, i.e. 4.2 kBit/s of net payload.
Assuming a 100% protocol overhead, exhausting a mature 100 MBit/s interface at the server side would imply that we are completely done in 2.5 hours.

Contributing Devices

Board	Device	Solvers	Clock	SE
VC707	XC7VX485T-2	325	250.0 MHz	812
KC705	XC7K325T-2	241	290.4 MHz	700
ML605	XC6VLX240T-1	125	200.0 MHz	250
DE4	EP4SGX230KF40C2	125	250.0 MHz	312
DNK7_F5_PCl_e	5 × XC7K325T-1	5 × 240	220.0 MHz	2640

SE (Solver Equivalent): one solver unit running at 100 MHz


Emergent Patterns: Solution Counts



(First 1000 *lexicographically ordered* subproblems)

Putting Queens in Carry Chains– No. 27 –

Computational Snapshot



State of Affairs

Currently:

- an average of 15 solutions per second in undisrupted operation is achieved, and
- 2.7% of the 2,024,110,796 subproblems are solved.

Ongoing Efforts:

- Use of local clock resources (BUFR, BUFH) is explored to squeeze out more performance.
- A GPU port is under development.

Thank you!

The whole implementation is available as open-source:
<https://github.com/preusser/q27>